Distributed Multi-Agent Decision Making Under Uncertain Communication

by

Cameron W. Pittman

M.A., Belmont University (2011)
B.A., Vanderbilt University (2009)

Submitted to the Department of Aeronautics and Astronautics
in partial fulfillment of the requirements for the degree of

Master of Science

at the

MASSACHUSETTS INSTITUTE OF TECHNOLOGY

June 2023

Author . . . .
Department of Aeronautics and Astronautics
May 1, 2023

Certified by. . . . .
Brian C. Williams
Professor of Aeronautics and Astronautics, MIT
Thesis Supervisor

Accepted by . . . .
Jonathan How
R.C Maclaurin Professor of Aeronautics and Astronautics, MIT
Chair, Graduate Program Committee

Distributed Multi-Agent Decision Making Under Uncertain Communication

by

Cameron W. Pittman

Abstract

This is an abstract.

Thesis Supervisor: Brian C. Williams
Title: Professor of Aeronautics and Astronautics, MIT

Acknowledgements

So long and thanks for all the fish!

THIS PAGE INTENTIONALLY LEFT BLANK

Contents

List of Figures

THIS PAGE INTENTIONALLY LEFT BLANK

List of source codes

THIS PAGE INTENTIONALLY LEFT BLANK

List of Tables

THIS PAGE INTENTIONALLY LEFT BLANK

Chapter 1

Introduction

This is the introduction.

1.1 Section Header

This is a section of text. As said in... [1] "Hey". So said [2] too that things are cool.

```
(format t "hello, world!")
```

Listing 1: This is a caption.

This is something I want to cite [3]

THIS PAGE INTENTIONALLY LEFT BLANK

Chapter 2

Problem Statement

There is a real-world need for coordinating multiple agents that are collaborating while facing uncertain inter-agent communication in uncertain environments. This thesis will focus on two motivating scenarios drawn from collaborative space operations, though this section will include a third example from the domain of military operations to further motivate the need for modeling uncertain delay in observations. The first scenario is based on an Artemis-like extravehicular activity (EVA), where an astronaut and a rover on the lunar surface are required to coordinate to share limited downlink bandwidth. The second scenario comes from the domain of satellite swarms, where many homogeneous agents coordinate operations with tight temporal constraints. The third scenario (which will not be modeled in experiments) is one where cooperative military agents are moving into and out of regions where immediate communication is purposefully halted in order to avoid detection by an enemy.

Delay temporal networks are essential to modeling space operations, where uncertainty is systemic to all aspects of planning, scheduling, and execution. As will be described in more detail for each motivating scenario below, uncertain observation delay is unavoidable due to imperfect communication infrastructure, uncertain timing with respect to knowledge transfer between agents, and finicky scientific equipment.

2.1 Extravehicular Activities as a Motivating Scenario

NASA is preparing to send astronauts back to the lunar surface as part of the Artemis program[1]. The Artemis astronauts will continue a longstanding tradition in the U.S. space program of performing EVAs [4]. Like the Apollo astronauts of the 1960s and 1970s, Artemis astronauts will embark on EVAs wherein crews don spacesuits, egress landers, and conduct scientific expeditions on the lunar surface. There, they will survey surface features, collect samples, and in general perform field geology

---

[1] `https://www.nasa.gov/specials/artemis/`

[5][7]. A small number of astronauts are Ph.D. geologists by trade, and NASA is training others in the principles of field geology up to a notional masters level of understanding [8]. NASA will support the lunar activities of these astronauts through a vast infrastructure of personnel on the ground, including teams of domain-relevant scientists. Together with flight controllers and engineers from various disciplines, a science team on the ground will provide real-time feedback to ensure that Artemis astronauts maximize the scientific return of their EVAs.

The relevant actors in an EVA include extravehicular (EV) crew members, who conduct all field activities outside the vehicles and habitats, and a ground-based Mission Control Center (MCC). Typically there are two EV crew members who often, but not always, work together to complete tasks. Life support imposes a temporal bound on the overall length of excursions. As an EVA progresses, EV crews consume four non-renewable resources, comprised of oxygen, battery power, water, and CO_2 scrubbers [9]. The duration of EVAs is limited by the consumable that is on track to be depleted first across the life support systems of both EV crew members, referred to as the limiting consumable.

Mission Control Center (MCC) is comprised of hundreds of flight controllers and engineers who monitor all aspects of EVAs. Its strict hierarchical structure ensures that all space-to-ground decisions pass through the Flight Director [10]. For the purposes of this example, we treat MCC as a single actor.

When astronauts perform field science, another actor comes into play, called the ground-based Science Backroom Team (SBT). The SBT is comprised of multidisciplinary scientists who help astronauts prioritize and select scientific sample targets online [5], [11]. The science team reports their priorities to MCC, who then passes them along to the crew. The SBT behaves as a separate actor with limited communication, in that their messages may only pass directly to MCC, not the crew.

For EVAs, delay is manifested across three distinct categories: signal transmission, human operational delays, and instrument processing. Each category presents a source of delay that varies with uncertainty. For signal transmission, delay is sourced from the speed of light between planetary bodies and infrastructural deficiencies. Communication infrastructure outside of low-Earth orbit, including satellites and planetary surface signal repeaters [12], is not robust, and as such unpredictable delays and signal dropouts will be common [13]. For human operational delays, note that the primary goal of Mission Control is keeping the crew safe [14]. As such, communications from the science team to the crew may be delayed or dropped because Mission Control needs to prioritize communications and actions related to crew health and safety at the expense of science [15], [16]. Lastly, for instrument processing, there is uncertainty in the temporal relationships between the activation of complex scientific instruments and the return of useful information NO_ITEM_DATA:Sehlke2019. Scientific packages may generate high and low bandwidth data products, the uplink of which will

be bottlenecked by limited bandwidth between space and ground, creating additional uncertainty between when a data product is ready and when it is available to the science team.

2.2 Motivating Scenario

A human and a rover are working together to perform scientific exploration on the lunar surface. The human is performing exploration of targets of opportunity that were identified during descent. The rover has a robotic arm, which it is using to perform sampling tasks collecting rocks in a predetermined location.

Both are in communication with mission control and scientists on Earth. Importantly, they share bandwidth on the relays and satellites used to transmit data between the Moon and Earth. CONOPS says only one agent may be using the bandwidth at a time (including uploads from Earth)

Both the human and the robot have their own schedules of events to perform and are aware of the other agent's events.

THIS PAGE INTENTIONALLY LEFT BLANK

Chapter 3

Approach

For addressing our problem statement, we envision that a high-level executive should take responsibility for managing task planning and execution with respect to temporal constraints. For this thesis, we chose to extend an existing high-level task and motion planner, *Kirk* [18]. Kirk is a complete, end-to-end executive in that it can take human-friendly problem specifications as input and send commands to hardware as output.

To clarify terminology in this thesis, the term *executive* refers to Kirk and its subsystems, while *agent* refers to the combination of an executive and the system it controls that interacts with the outside world, e.g. robotic hardware.

At a high-level, Kirk works by first taking a description of the problem domain as written by domain experts, which should include the constraints, agent dynamics, environment, and starting and goal states of the problem at hand. Kirk then generates and checks plans using an optimal satisfiability (OpSAT) solver [19], elaborates plans to sub-executives when it encounters constraints and goals it cannot plan against directly, and eventually dispatches event schedules and motion plans to hardware. For the purpose of this thesis, we primarily focus on Kirk's capability to dispatch events, though fully accounting for uncertain communication in a real agent requires that all of Kirk's aforementioned capabilities are updated.

Some aspects of Kirk were already well-suited for coordinating multiple agents under observation delay, others were not. Specifically, our approach required research contributions in three key areas, which were then implemented in Kirk:

1. *Modeling and Controllability*: prior to execution, we must be able to model communication delay separate from temporal constraints, as well as guarantee that all temporal constraints can be satisfied

2. *Scheduling*: during execution, executives must be able to dynamically schedule and dispatch events respecting temporal constraints in spite of observation delay

3. *Coordination*: during execution, peer executives must be able to share event assignments and observations

We include an auxiliary fourth area for contributions as an engineering requirement.

4. *Robustness*: Kirk must be easy to run, debug, and integrate with existing autonomous systems

Our approach to each research focus will be described below.

3.1   Modeling and Controllability

We take a model-based approach to deploying autonomous systems, that is, prior to a mission, we envision that engineers and domain experts work together to model the system at hand, then during the mission (though not necessarily online), the autonomous system then takes the models as input and decides how to act as output. There are three core challenges with modeling - the first being that we need formalisms that can be ingested by our algorithms and be used to guarantee the safe execution. In other words, we need a data type to represent the phenomenon over which we want the algorithms comprising our system to reason. Next, the chosen formalism must allow us to guarantee the satisfiability of the system in that the autonomous system must be able to act in a safe manner respecting all constraints to go from the starting state to the goal state. Finally, the third challenge is that we need a human-friendly form of said formalisms such that human domain experts, who are unlikely to also be experts in autonomy, can still model their domains accurately enough such that the desired safe behavior is output by the autonomous system. We address both challenges in our approach to modeling.

States and constraints can take on arbitrary forms, and how they are modeled depends entirely on the problem domain. Classical planning problems use boolean predicates and actions to model the world (e.g. STRIPS planning problems [20]). Scheduling problems involving time constraints will have continuous temporal bounds between discrete timepoints (e.g. in the form of temporal constraint graphs [21]). Other scenarios where motion planning is the focus will likely be modeled with vectors of continuous values in $\mathbb{R}$ (e.g. often representing convex regions as in the case of the *Magellan* planner [22]). Hybrid domains combine states and constraints with mixed continuous and discrete values (e.g. using mixed-integer linear programs as demonstrated by Chen et al. [23]).

Given this thesis' emphasis on temporal scheduling, we choose to focus entirely on formalisms where states and constraints are temporal in nature. The starting state of the system is, by definition, one where time is set to 0 seconds, $t = 0$, and no events have been executed (i.e. no event assignments have been made). We then define controlled and uncontrolled set-bounded constraints between events. The goal state is one where times have been assigned to each controllable event such that all constraints are satisfied. To do so, we build our formalisms representing temporal constraints

with set-bounded observation delay on top of simple temporal networks with uncertainty (STNUs) [24]. A brief explanation of our modeling strategy for temporal constraints with observation delay follows in Section 3.1.1, though we will elaborate on temporal reasoning and our chosen formalisms for it in much more detail in Chapter 4.

With a modeling formalism in hand, the second key challenge is to use the formalism to guarantee a property known as *controllability*, or that all controllable temporal constraints can be satisfied given the existing uncertainty in the STNU. There already exist a number of strategies for checking the controllability of STNUs. Examples of different strategies include the canonical work by Morris, Muscettola, and Vidal in checking for semi-reducible negative cycles (SRNCs) [25][28], as well as more exotic approaches like reframing controllability as a Satisfiable Modulo Theory (SMT) problem [29]. In our approach to controllability under observation uncertainty, we build on top of checks for SRNCs as will be shown in 4.3.

For the third challenge, we choose to extend the Reactive Model-Based Programming Language (RMPL) [30], which provides to domain experts a means for describing the constraints and goal states of their domain without requiring additional expert knowledge in autonomy. With RMPL, a human planner is capable of building control programs describing the constraints, agents, and states of the problem domain in a way that is human-readable yet highly programmable, and is independent of the underlying algorithms used by the autonomous system. As will be explained in Section 3.1.2 below, our approach was to add the ability for planners to model observation delay alongside temporal constraints in RMPL.

3.1.1 Modeling Uncertain Observation Delay in STNUs

In the case of observation delay, our model dictates that we reason over two time intervals. The first time interval represents the true length of time between two events, while the second interval represents the length of time between when an event occurs and when an executive observes the event. For ensuring that an executive takes safe actions in an uncertain environment, we assume worst-case scenario with respect to information gain. Our approach to modeling uncertain observation delay in STNUs is as follows.

1. The duration of time between two events is represented as a set-bounded interval

2. The duration of time between an event and its observation (observation delay) is represented as a set-bounded interval

3. Timestamps in event observations are ignored

4. The true duration of observation delay is not guaranteed to be learned

The first point comes directly from the STNU formalism (see Section 4.1). The second point allows for uncertainty in the amount of observation delay, e.g. in an uncertain environment, we could model observation delay for a given event as, say, $[1, \infty]$, meaning an observation of an event could arrive one second after it occurs, or never arrive, or arrive at some arbitrary time, t, $1 < t < \infty$ later. The third point comes from assuming worst-case scenario and prevents us from "cheating" in our scheduling algorithm. For instance, imagine two agents coordinating. If agents passed timestamp information along with events to one another, they must also be able to synchronize their clocks, potentially to an arbitrary degree of precision. The challenge of synchronizing clocks between agents is outside the scope of this thesis and may not always be possible. As such, executives only trust their own clocks. Rather than backfill potentially erroneous times for event assignments as reported by exogenous sources, the executive we envision in this thesis records times that are internally consistent with its own clock. Doing so guarantees that the actions the executive takes as a result of temporal reasoning are consistent with its model.

The fourth point, that we are not guaranteed to learn event assignments, is a result of the first three. It stands to reason that an event observation is a function of the true assignment of an event and its observation delay. If there is uncertainty in both the event assignment and delay, then we have one equation with two unknowns. Thus, the term "uncertain" in uncertain observation delay means that we are forced to reason with deciding when to act even when we are not guaranteed to learn the true times assigned to events.

We formalize event observations and observation delay in Section 4.3.

3.1.2 Modeling Observation Delay in RMPL

RMPL is a key component of Kirk. This section steps through example RMPL control programs to describe their features and our modeling choices. The purpose of this section is three-fold:

1. A short walkthrough of the language is required in order to explain this thesis' contributions because an updated RMPL description in any form (e.g. manual, publication, or tutorial) has not been publicly released since 2003 [18]

2. We must describe the modeling choices of RMPL in sufficient detail to make concrete our approach to modeling temporal constraints in human-readble form

3. The above is used to demonstrate that modeling uncertain communication delay can be naturally modeled in RMPL

This section is not meant to be a complete documentation of RMPL, rather our goal is to motivate the strength of RMPL as a modeling language for human planners describing autonomous systems with observation uncertainty.

RMPL has undergone a number of rewrites since its inception, and is currently being developed as a superset of the Common Lisp language using the Metaobject Protocol [31]. The goal is that a human should have a comfortable means for accurately modeling sufficient detail about the problem domain such that an executive can perform model-based reasoning to decide how to act.

RMPL and Kirk can be used to achieve a number of different goals. These include but are not limited to temporal scheduling, classical planning, hybrid planning. For this thesis, we focus on temporal scheduling and the ability for a human to write *control programs*, or composable constraints and goals.

For this thesis, we take the assumption that each Kirk executive is responsible for a single agent. We also ignore vehicle dynamics given this thesis' focus on contributions to temporal scheduling. However, RMPL is more flexible and allows multi-agent planning and motion planning using vehicle dynamics, which will be briefly described in Section 3.1.3.

An example of an RMPL control program for a single-agent without agent dynamics follows in Listing 2.

```lisp
;; NOTE: we omitted Lisp package definitions here for simplicity's sake

(define-control-program eat-breakfast ()
  (declare (primitive)
           (duration (simple :lower-bound 15 :upper-bound 20))))

(define-control-program bike-to-lecture ()
  (declare (primitive)
           (duration (simple :lower-bound 15 :upper-bound 20))))

(define-control-program main ()
  (with-temporal-constraint (simple-temporal :upper-bound 40)
    (sequence (:slack nil)
              (eat-breakfast)
              (bike-to-lecture))))
```

Listing 2: A sample control program composed of three constraints. `eat-breakfast` and `bike-to-lecture` designate controllable constraints, while the `main` control program enforces that the constraints are satisfied in series.

Looking past the parentheses, we can see different options for defining temporal constraints. For example, the `(duration (simple ...))` form is used to define a set-bounded temporal constraint between a `:lower-bound` and an `:upper-bound`. The `main` control program uses a different form, `(with-temporal-constraint ...)` to place an `:upper-bound` on the overall deadline for scheduling all events in the control program.

The example control programs in Listing 2 are defined without agents in that there is an assumption that the Kirk instance that executes this control program must know what the semantics of `eat-breakfast` and `bike-to-lecture` mean and how to execute them.

It could also be the case that Kirk is simply being used to produce a schedule of events offline that will be handed to an agent that knows how to execute them. As an example, perhaps a student wants some help planning their morning, so they write an RMPL control program with constraints representing everything they need to do between waking up and going to lecture, as seen in the more complex control program in Listing 3. The student could ask Kirk to produce a schedule of events that satisfies all the temporal constraints in this RMPL control program, which they would then use to plan their morning routine. See the resulting schedule produced by Kirk in Table 3.1. (Note that while normally times in RMPL are represented in seconds, we use minutes in Listing 3 and Table 3.1 for simplicity's sake.)

Table 3.1: The schedule produced by Kirk's scheduler for the student's routine before lecture as modeled in Listing 3. Note: Kirk's output has been cleaned for readability purposes.

Event	Time (min)
`START`	0
Start `shower`	1
End `shower`	6
Start `review-scheduling-notes`	6
Start `eat-breakfast`	6
End `review-scheduling-notes`	16
Start `review-planning-notes`	16
End `eat-breakfast`	21
End `review-planning-notes`	26
Start `pack-bag`	26
End `pack-bag`	31
Start `bike-to-lecture`	32
End `bike-to-lecture`	46
`END`	46

Listing 3 introduces the notion of control programs that are allowed to be executed simultaneously, as modeled with the `(parallel ...)` form found in the `main` control program on line 48.

Kirk is able to simulate the RMPL script in Listing 3 and produce a schedule because there were no uncontrollable constraints, that is, all control programs are under the agent's control. Say we replaced `bike-to-lecture` with `drive-to-lecture`. Due to traffic conditions, driving presents in an uncontrollable constraint. RMPL allows us to model uncontrollable constraints as in Listing 4.

The addition of `:contingent t` to the `(duration ...)` form tells Kirk that it does not have control over when the end of `drive-to-lecture` is scheduled, rather, Nature (i.e. traffic conditions) chooses a time. Despite the lack of control over `drive-to-lecture`, we do know the drive should take between 15 and 20 minutes, hence our model includes `:lower-bound 15` and `:upper-bound 20`.

With uncontrollable constraints in a control program, we are no longer guaranteed to be able to produce a schedule offline as we show in Table 3.1. Instead, as time passes, we may only choose to

```
1   ;; This file lives in the thesis code repo at:
2   ;;       kirk-v2/examples/morning-lecture/script.rmpl
3   ;;
4   ;; To execute this RMPL control program as-is and generate a schedule, go to the root
5   ;; of the thesis code repo and run the following command:
6   ;;
7   ;; kirk run kirk-v2/examples/morning-lecture/script.rmpl \
8   ;;       -P morning-lecture \
9   ;;       --simulate
10
11  (rmpl/lang:defpackage #:morning-lecture)
12
13  (in-package #:morning-lecture)
14
15  (define-control-program shower ()
16    (declare (primitive)
17             (duration (simple :lower-bound 5 :upper-bound 10))))
18
19  (define-control-program eat-breakfast ()
20    (declare (primitive)
21             (duration (simple :lower-bound 15 :upper-bound 20))))
22
23  (define-control-program review-scheduling-notes ()
24    (declare (primitive)
25             (duration (simple :lower-bound 10 :upper-bound 15))))
26
27  (define-control-program review-planning-notes ()
28    (declare (primitive)
29             (duration (simple :lower-bound 10 :upper-bound 15))))
30
31  (define-control-program pack-bag ()
32    (declare (primitive)
33             (duration (simple :lower-bound 5 :upper-bound 6))))
34
35  (define-control-program bike-to-lecture ()
36    (declare (primitive)
37             (duration (simple :lower-bound 15 :upper-bound 20))))
38
39  (define-control-program review-notes ()
40    (sequence (:slack t)
41      (review-scheduling-notes)
42      (review-planning-notes)))
43
44  (define-control-program main ()
45    (with-temporal-constraint (simple-temporal :upper-bound 60)
46      (sequence (:slack t)
47        (shower)
48        (parallel (:slack t)
49          (eat-breakfast)
50          (review-notes))
51        (pack-bag)
52        (bike-to-lecture))))
```

Listing 3: A student's morning routine preparing for lecture as modeled in RMPL. This is a complete RMPL program that includes the required Lisp package definitions to run in Kirk.

```
(define-control-program drive-to-lecture ()
  (declare (primitive)
           (duration (simple :lower-bound 15 :upper-bound 20)
                     :contingent t)))
```

Listing 4: An uncontrollable, or contingent, temporal constraint in a control program.

schedule controllable events based on the *partial history* of contingent event assignments so far, or, in other words, perform *dynamic scheduling*. Thus, we can no longer simulate a schedule with Kirk. We must connect Kirk to a source for receiving contingent event assignments in order to make valid controllable event assignments. Our approach to dynamic scheduling is the focus of Section 3.2.

As a contribution of this thesis, our existing approach to specifying durations in RMPL was expanded to model observation delay. An example follows in Listing 5 modeling a sample collection control program with observation delay.

```
(define-control-program collect-science-sample ()
  (declare (primitive)
           (duration (simple :lower-bound 15 :upper-bound 30
                             :min-observation-delay 5
                             :max-observation-delay 15)
                     :contingent t)))
```

Listing 5: An RMPL control program describing a science data collection task with observation delay.

We can see in Listing 5 that representing set-bounded observation delay is a simple as adding `:min-` and `:max-observation-delay` to the `(duration (simple ...)  :contingent t)` form. In full, this control program represents an uncontrollable constraint with a contingent event that Nature will schedule $[15, 30]$ time units after sample collection begins. The executive will then wait an additional $[5, 15]$ time units before learning that `collect-science-sample` has been scheduled. As will be described in much greater detail in Section 4.3, the executive will only learn *that* the contingent event occurred - is not guaranteed to learn where in $[15, 30]$ the contingent event was assigned, nor will it know how much observation delay was incurred.

3.1.3 Explicitly Modeling Agents in RMPL

This section is included to expand on the features of RMPL, though note that none of these features are required for controlling distributed agents, and were not a part of the experiments for this research.

If we wanted to specify agents in a multi-agent control program, or if we wanted to take vehicle dynamics into account, RMPL gives us a means for using the Common Lisp Object System (CLOS) for defining agents, agent dynamics, and the control programs agents may execute.

An example RMPL control program with an agent is provided in Listing 6 for completeness sake from the domain of underwater robotics.

In Listing 6, `glider` refers to a low-powered autonomous underwater vehicle that prefers to traverse by following ocean currents using a buoyancy engine.[1] We see that we model a `glider` agent and its properties using standard CLOS. The `move` control program then takes a `glider` and a `location` as arguments. The (`requires ...`) form is equivalent to the preconditions of a durative action in a PDDL 2.1 [32] domain. Likewise, the (`effect ...`) form is equivalent to PDDL effects. Finally, as we saw before, the durative action also includes a temporal constraint in its (`duration ...`) form.

Kirk is able to take RMPL as input to perform classical planning, though further discussion of it falls outside the scope of this thesis.

3.2   Scheduling Temporal Events

The bulk of the technical chapters of this thesis, namely Chapters 4 and 5, describe the algorithmic insights behind the *delay scheduler*. The delay scheduler dispatches controllable events online for dynamically controllable STNUs while reasoning over observation delay in the uncontrollable events it receives. There were two key contributions that enabled the delay scheduler.

Reasoning over the controllability of STNUs with with variable-observation delay had been demonstrated to be possible in prior work [33], though a valid execution strategy was never defined. For our first contribution, we define an execution strategy for variable-delay controllable STNUs and prove its validity.

Likewise, dynamic schedulers have been established for dispatching events from STNUs, e.g. FAST-EX [34]. For our second contribution, we defined a novel delay scheduler built on FAST-EX capable of applying the execution strategy defined in our first contribution.

We elaborate further on our approach to each contribution below.

3.2.1   Defining a Valid Execution Strategy for STNUs with Variable Observation Delay

We cannot execute an STNU without first demonstrating that it is controllable. Our approach to checking the controllability of STNUs with observation delay is to apply Bhargava's Variable-Delay Controllability checker (VDC) [1]. VDC is a procedure that takes place in two stages and is $O(N^5)$ in the number of events. In the first stage, we transform the STNU with variable observation delay to one with fixed observation delay in $O(N^2)$. In the second stage, we check the controllability of

---

```lisp
;; This code is a snippet from a file in the thesis code repo found at:
;;      kirk-v2/examples/glider/script.rmpl

(defclass glider ()
  ((id
    :initarg :id
    :finalp t
    :type integer
    :reader id
    :documentation
    "The ID of this glider.")
   (deployed-p
    :initform nil
    :type boolean
    :accessor deployed-p
    :documentaiton
    "A boolean stating if the glider is deployed at any point in time.")
   (destination
    :initform nil
    :type (member nil "start" "end" "science-1" "science-2")
    :accessor destination
    :documentation
    "The location to which the glider is currently heading, or NIL if it is not
    in transit.")
   (location
    :initarg :location
    :initform "start"
    :type (member nil "start" "end" "science-1" "science-2")
    :accessor location
    :documentation
    "The location where the glider is currently located, or NIL if it is not at
    a location (in transit).")))

(define-control-program move (glider to)
  (declare (primitive)
          (requires (and
                      (over :all (= (destination glider) to))))
          (effect (and
                      (at :start (= (destination glider) to))
                      (at :start (= (location glider) nil))
                      (at :end (= (destination glider) nil))
                      (at :end (= (location glider) to))))
          (duration (simple :lower-bound 10 :upper-bound 20))))
```

Listing 6: A snippet of an RMPL script that defines an agent and classical planning predicates and effects of a control program.

the fixed-delay STNU using Bhargava's fixed-delay controllability checker (FDC) [33], [35], which is modified from Morris' $O(N^3)$ dynamic controllability check [28] such that it accounts for fixed observation delay in contingent links.

In short, the first stage process is built around the idea of modeling a worst-case scenario with respect to receiving observations. The resulting fixed-delay STNU reflects a situation where the executive learns as little as possible about the contingent events. If the fixed-delay STNU with minimal information is controllable, then so too must any situation be controllable when we learn more information.

Our first contribution is that we define an execution strategy for variable-delay STNUs, wherein we dispatch events according to the *dispatchable form* of the *fixed-delay* STNU, while respecting the constraints modeled in the *variable-delay* STNU. Existing controllability checks, like FDC, and execution strategies for STNUs, like FAST-EX, depend on a dispatchable form, or a distance graph representation of the STNU. The key challenge in defining an execution strategy for a variable-delay STNU is that unlike vanilla STNUs and fixed-delay STNUs, there is no dispatchable form for variable-delay STNUs. Hence why the VDC check first transforms the variable-delay STNU to a fixed-delay form. In Chapter 5, we formally define the execution strategy for variable-delay STNUs and prove its validity.

3.2.2   Online Dispatching for STNUs with Variable Observation Delay

We chose to build the delay scheduler as a modified variant of Hunsberger's FAST-EX [34] because FAST-EX appears to be the fastest dynamic scheduler published to date. To improve the online performance of dynamic scheduling, Hunsberger's insight was to reduce the space of the dispatchable form by removing edges as events are executed. We can do so because we iteratively update the distances to and from the remaining events by performing Dijkstra's Single Sink and Single Source Shortest Paths algorithms to and from the zero point (start event) of the distance graph.

FAST-EX maps partial histories, or schedules of events up to the current time, to Real-Time Execution Decisions (RTEDs). RTEDs contain a list of events to be executed and a time (that could be from now to point in the future) to execute them.

The delay scheduler differs from FAST-EX in the way it records partial histories and how it generates RTEDs.

3.3   Coordination

Coordination between Kirk agents is a completely novel contribution of this thesis.

3.4 Robustness

We improve the delay scheduler by differentiating real and `noop` controllable events. . .

Chapter 4

Modeling Temporal Networks

Our aim is to architect a pipeline that takes a temporal network as input and produces temporally consistent actions as output. This Section lays the groundwork for the delay scheduler by first describing our chosen models for temporal constraints and scheduling algorithms. We then outline the architecture of the delay scheduler, before finally providing the necessary background on our choice for modeling observation delay.

TODO ... take the approach of defining temporal networks, and updating our defns as we work towards STNUs

4.1 Temporal Networks

Temporal networks form the backbone of our architecture for temporal reasoning under observation delay. Simple Temporal Networks (STNs) offer the basic building blocks for most expressive temporal network formalisms [21]. An STN is composed of a set of variables and a set of binary constraints, each of which limits the difference between a pair of these variables; for example, $B - A \in [10, 20]$. Each variable denotes a distinguished point in time, called an *event*. Constraints over events are binary *temporal constraints* that limit their temporal difference; for example, the aforementioned constraint specifies that event A must happen between 10 and 20 minutes before event B.

Definition 1. STN [21]

An *STN* is a pair $\langle X, R \rangle$, where:

- X is a set of variables, called events, each with a domain of the reals $\mathbb{R}$, and

- R is a set of simple temporal constraints. Each constraint $\langle x_r, y_r, l_r, u_r \rangle$ has scope $\{x_r, y_r\} \subseteq X$ and relation $x_r - y_r \in [l_r, u_r]$.

Definition 2. Schedule [25]

A *schedule*, ξ, is a mapping of events to times, $\xi : X \to \mathbb{R}$.

An STN is used to frame scheduling problems. A schedule is feasible if it satisfies each constraint in R. We use the notation $\xi(x)$ to represent a mapping from an event, x, to a time, $x \to \mathbb{R}$, in the schedule. A schedule is *complete* if all $x \in X$ are assigned times in ξ. An STN is *consistent* if it has at least one feasible schedule that assigns all events in X.

An STN is consistent if and only if there is no negative cycle in its equivalent distance graph [21]. Let n be the number of events in a temporal network and m to be the number of constraints. Then consistency of an STN can be checked in $O(mn)$ time using the Bellman-Ford Algorithm to check for negative cycles.

While an STN is useful for modeling problems in which an agent can control the exact time of all events, it does not let us model actions whose durations are uncertain. A Simple Temporal Network with Uncertainty (STNU) is an extension to an STN that allows us to model these types of uncertain actions [24].

Definition 3. STNU [24]

An *STNU S* is a quadruple $\langle X_e, X_c, R_r, R_c \rangle$, where:

- X_e is the set of executable events with domain $\mathbb{R}$,

- X_c is the set of contingent events with domain $\mathbb{R}$,

- R_r is the set of requirement constraints of the form $l_r \leq x_r - y_r \leq u_r$, where $x_r, y_r \in X_c \cup X_e$ and $l_r, u_r \in \mathbb{R}$, and

- R_c is the set of contingent constraints of the form $0 \leq l_r \leq c_r - e_r \leq u_r$, where $c_r \in X_c$, $e_r \in X_e$ and $l_r, u_r \in \mathbb{R}$.

An STNU divides its events into executable and contingent events and divides its constraints into requirement and contingent constraints. The times of executable events are under the control of an agent, and assigned by its scheduler. STNU executable events are equivalent to events in an STN. Contingent events are controlled by nature. Contingent constraints model the temporal consequences of uncertain actions and are enforced by nature. They relate a starting executable event and an ending contingent event. To ensure causality, the lower-bound of a contingent constraint is required to be non-negative; hence, the end event of the constraint follows its start event. Contingent constraints are not allowed to be immediately followed by additional contingent constraints. Requirement constraints specify constraints that the scheduler needs to satisfy and may relate any pair of events. An STNU requirement constraint is equivalent to an STN constraint.

To clarify terminology, we sometimes refer to contingent constraints as contingent links and requirement constraints as requirement links. When we discuss contingent constraint or contingent

link duration, we refer to the amount of time that actually elapses between a contingent link's starting executable event and its ending contingent event. We sometimes refer to STNUs as defined in Defintion 3 as *vanilla* STNUs (in contrast to the many "flavors" of STNUs, namely the variants with fixed and variable observation delay functions as will be defined below).

With STNs, our goal is to construct a consistent schedule for all events such that all constraints are satisfied. In STNUs, however, contingent events cannot be scheduled directly. Instead, we are interested in determining whether there is a *controllable* policy that guarantees that a schedule can be constructed such that all constraints are satisfied despite how uncertainty is resolved.

Definition 4. Situations [24]

For an STNU S with k contingent constraints $\langle e_1, c_1, l_1, u_1 \rangle, \cdots, \langle e_k, c_k, l_k, u_k \rangle$, each *situation*, ω, represents a possible set of values for all links in S, $\omega = (\omega_1, \cdots, \omega_k) \in \Omega$. The *space of situations* for S, Ω, is $\Omega = [e_1, c_1] \times \cdots \times [e_k, c_k]$.

Each *situation* in the *space of situations*, $\omega \in \Omega$, represents a different assignment of contingent links in the schedule [24]. We may represent the situation for a specific constraint as ω_i for the i-th constraint in S, or $\omega(E)$ for contingent event E. Situations are sets of intervals. To examine spaces of situations, we can make the following comparisons.

Definition 5. Comparisons of Spaces of Situations

Given two spaces of situations, Ω_1 and Ω_2, with contingent link j, $1 \leq j \leq k$,

- $\Omega_1 = \Omega_2$ if and only if $\omega_1 = \omega_2 \forall \omega_1 \in \Omega_1 \forall \omega_2 \in \Omega_2$

- $\Omega_1 \subset \Omega_2$ if situation j in Ω_1 is a subset of situation j in Ω_2, $\omega_{1j} \subset \omega_{2j}$, and all other situations are equivalent

- $\Omega_1 \subset \Omega_2$ if Ω_1 omits contingent link j, e.g. $\Omega_1 = \prod_{\substack{i=1 \\ i \neq j}}^{k} [e_{1i}, c_{1i}]$, and all other situations are equal.

Situations may be applied to STNUs.

Definition 6. Projection [24], [25]

A *Projection* is an application of a situation, ω, on an STNU S, which collapses the durations of contingent links to specific durations resulting in an STN.

A *projection* is an STN that is the result of applying a situation to an STNU, and thus the contingent links have reduced from uncertain ranges to specific durations [24], [25].

Definition 7. Execution Strategy

An *execution strategy*, $\mathcal{S}$, is a mapping of situations to schedules, $\mathcal{S} : \Omega \rightarrow \Xi$.

An *execution strategy* then naturally maps a specific resolution of the uncertainty of the contingent constraints to a set of assignments for the events of an STNU. For an STNU, time monotonically increases and we only observe *activated* contingent events, or those contingent events at the tail of a contingent link whose free event predecessor has been executed. As such, we modify our definition of ξ.

Definition 8. Partial Schedule

A *partial schedule*, ξ, is a mapping from a proper subset of events in X, X', to times, $\xi : X' \to \mathbb{R}$.

As a proper subset, ξ represents an assignment of events *so far* during the execution of an STNU. From here on, ξ refers to a partial schedule.

In the world of STNU literature, there are many forms of controllability that represent the ability of a scheduler to construct execution strategies that satisfy constraints under different conditions [24]. Three forms of controllability, *strong*, *weak*, and *dynamic* are studied most often, though in practice we omit weak controllability from our analysis. As we will see below, variable-delay controllability will unify strong and dynamic controllability into a single theory. A temporal network is *strongly controllable* (or exhibits strong controllability), if there exists a complete schedule that will satisfy all constraints for all projections of the STNU. A temporal network exhibits dynamic controllability if an execution strategy, $\mathcal{S}$, exists for a given partial schedule, ξ.

4.2 Fixed-Delay Controllability

Under fixed-delay controllability [33], we consider the problem of scheduling execution decisions when the assignment of values to contingent events is learned (if at all) after some time has passed from the initial assignment. We prefer this model because it is flexible enough to model most forms of event observation we would expect to see in a real-time execution context. Fixed-delay controllability uses a *fixed-delay function* to encode the delay between when an event occurs and when it is observed by a scheduling agent. We sometimes refer to an STNU with an associated fixed-delay function as a *fixed-delay STNU*.

Definition 9. Fixed-Delay Function [33]

A *fixed-delay function*, $\gamma : X_c \to \mathbb{R}^+ \cup \{\infty\}$, maps a contingent event to the amount of time that passes between when the event occurs and when its value is observed.

As a matter of convention, we use $A \xrightarrow{[l,u]} B$ to represent requirement links between events A and B and use $A \overset{[l,u]}{\Longrightarrow} E$ to represent contingent links between A and E. When we refer to the fixed-delay function associated with a contingent event E of some contingent constraint $A \overset{[l,u]}{\Longrightarrow} E$, we use the notation $\gamma(E)$, or equivalently, γ_E. Without instantaneous observation of contingent events, we must clarify the relationship between when an event occurs and when it is *observed*.

Definition 10. Contingent Event Observation

Observations, ψ, are a mapping from contingent events to times when the agent receives events, $\psi : E \to \mathbb{R}$, based on the relationship, $\psi(E) = \xi(E) + \gamma(E)$.

Note that $\xi(E)$ is now indirectly learned through the relationship,

$$\xi(E) = \psi(E) - \gamma(E)$$

We also present a revised definition of situations, Ω_f, to reflect the impact of the delay function on event observations.

Definition 11. Fixed-Delay Situations

For an STNU S with k contingent constraints $\langle e_1, c_1, l_1, u_1 \rangle, \cdots, \langle e_k, c_k, l_k, u_k \rangle$ and fixed-delay function γ, each *fixed-delay situation*, ω_f, represents a possible set of *observed* values for all links in S, $\omega_f = (\omega_{f1}, \cdots, \omega_{fk})$. The {space of situations} for S, Ω_f, is $\Omega_f = [e_1, c_1] + [\gamma_1, \gamma_1] \times \cdots \times [e_k, c_k] + [\gamma_k, \gamma_k]$.

To emphasize that the *observed* value for an event is not the same as $\xi(E)$, we also use the term *observation space* as a synonym for the space of situations.

With the semantics of delayed observations in hand, we can define what it means for a fixed-delay STNU to be controllable.

Definition 12. Fixed-Delay Controllability [33]

An STNU S is *fixed-delay controllable* with respect to a delay function, γ, if and only if for the space of situations, Ω_f, there exists an $\mathcal{S}$ that will construct a satisfying schedule for requirement constraints during execution, ξ.

Importantly, fixed-delay controllability (FDC) generalizes the two concepts of controllability that are central to STNUs, strong and dynamic controllability. In particular, by using a fixed-delay function where we observe all events instantaneously, checking fixed-delay controllability reduces to checking *dynamic controllability*. Similarly, a fixed-delay function that specifies we never observe any contingent event corresponds to checking *strong controllability* [24].

To determine whether an STNU is fixed-delay controllable, we determine whether there exists a *valid* execution strategy for it.

Definition 13. Valid Execution Strategy

A *valid* $\mathcal{S}$ is one that enforces that, for any $\omega_f \in \Omega_f$, while receiving information about the durations after a fixed delay, the outputted decision respects all existing temporal constraints and ensures the existence of a subsequent valid execution strategy following that action.

As is the case for a vanilla STNU, evaluating whether a valid $\mathcal{S}$ exists for a fixed-delay STNU reduces to checking for the presence of a *semi-reducible negative cycle* in a *labeled distance graph*

derived from the fixed-delay STNU [27]. The key insight for checking fixed-delay controllability is the inclusion of γ in the constraint generation rules for building the labeled distance graph [**?**].

The labeled distance graph corresponds to the constraints of the STNU with each unlabeled edge from A to B with weight w (denoted $A \xrightarrow{w} B$) representing the inequality $x_B - x_A \leq w$. Labeled edges represent conditional constraints that apply depending on the realized value of contingent links in the graph. For example, a lower-case labeled edge from A to B with weight w and lower-case label c (denoted $A \xrightarrow{c:w} B$) indicates that $x_B - x_A \leq w$ whenever the contingent link ending at C takes on its lowest possible value. An upper-case labeled edge from A to B with weight w and upper-case label C (denoted $A \xrightarrow{C:w} B$) indicates that $x_B - x_A \leq w$ whenever the contingent link ending at C takes on its highest possible value. Given a labeled distance graph, there are several valid derivations we can apply to generate additional edges (see Table 4.1). If it is possible to derive a negative cycle that is free of lower-case edges, then the STNU has a *semi-reducible negative cycle* and the STNU is not controllable.

Note that with fixed-delay controllability, the lower-case and cross-case rules are modified from the Morris and Muscettola [26], accounting for γ. More specifically, we address the case where observation delay makes it impossible to receive information about a contingent event before its immediate successor. More detail can be found in [36].

Edge Generation Rules			
	Input edges	Conditions	Output edge
No-Case Rule	$A \xrightarrow{u} B,\ B \xrightarrow{v} C$	N/A	$A \xrightarrow{u+v} C$
Upper-Case Rule	$A \xrightarrow{u} D,\ D \xrightarrow{C:v} B$	N/A	$A \xrightarrow{C:u+v} B$
Lower-Case Rule	$A \xrightarrow{c:x} C,\ C \xrightarrow{w} D$	$w < \gamma(C),\ C \neq D$	$A \xrightarrow{x+w} D$
Cross-Case Rule	$A \xrightarrow{c:x} C,\ C \xrightarrow{B:w} D$	$w < \gamma(C),\ B \neq C \neq D$	$A \xrightarrow{B:x+w} D$
Label Removal Rule	$B \xrightarrow{C:u} A,\ A \xRightarrow{[x,y]} C$	$u > -x$	$B \xrightarrow{u} A$

Table 4.1: Edge generation rules for a labeled distance graph from [36].

4.3 Variable-Delay Controllability

While fixed-delay controllability is quite expressive, its fundamental limitation is that it cannot model variability in delay. If uncertainty in delay is added to the model, then the controllability problem requires the addition of reasoning over uncertainty in both the timing of contingent events and their subsequent observations. Notably, with variable observation delay, we no longer guaranteed to learn the true time when contingent events were assigned.

We now introduce this model in terms of definitions for a *variable-delay function* and *variable-delay controllability* [1] as applied to *variable-delay STNUs*. Since variable-delay semantics generalizes the notion of fixed-delay, as a matter of convenience, we also use the simplified term *delay*

STNUs to refer to STNUs with variable observation delay.

Definition 14. Variable-Delay Function

A *variable-delay function*, $\bar{\gamma} : X_c \to (\mathbb{R}^+ \cup \{\infty\}) \times (\mathbb{R}^+ \cup \{\infty\})$, maps a contingent event to an interval $[a, b]$, where $a \leq b$. The interval bounds the time that passes after $\xi(x_c)$ before that value is observed to be assigned. No prior knowledge is assumed about the distribution associated with this interval.

Importantly, this model does not assume that an executing agent infers *when* a contingent event was executed, but instead infers *that* the event was executed. Like contingent constraints, the resolved value of $\bar{\gamma}$ will be selected by Nature during execution. Thus, when an agent receives an observation is a function of the independent resolutions of the contingent link and $\bar{\gamma}$.

By convention, we use $\bar{\gamma}^-(x_c)$ and $\bar{\gamma}^+(x_c)$ to represent the lower-bound and upper-bound, respectively, of the range representing the possible delay in observation. For any fixed-delay function γ, we can produce a corresponding variable-delay function $\bar{\gamma}$ where $\bar{\gamma}^+(x_c) = \bar{\gamma}^-(x_c) = \gamma(x_c)$.

During execution, the *observation projection*, Γ, represents the resolution of observation delay.

Definition 15. Observation Projection

The *observation projection* Γ is a mapping from a contingent event to a fixed observation delay, $\Gamma : X_c \to \mathbb{R} \in [\bar{\gamma}^-(X_c), \bar{\gamma}^+(X_c)]$.

Much like how a projection collapses a vanilla STNU to an STN, the observed projection collapses a contingent link with variable-observation delay to one with fixed-observation delay. However, unlike the projection of an STNU, the observation projection is not guaranteed to be learned. We update our definitions of ψ, ξ, and Ω accordingly.

Definition 16. Contingent Event Observation

Contingent event observations, ψ, are a mapping from contingent events to times when the agent receives events, $\psi : E \to \mathbb{R}$, based on the relationship, $\psi(E) = \xi(E) + \Gamma(E)$.

Determining the precise schedule of a contingent event, ξ, is complicated by an interval bounded Γ. We must use interval-bounded contingent event assignments instead.

Definition 17. Schedule

A *schedule*, ξ, when applied to contingent events, is a mapping of events to interval-bounded times, $\xi : X_c \to (\mathbb{R}^+ \cup \{\infty\}) \times (\mathbb{R}^{++} \cup \{\infty\})$, where, for any contingent event x_c, $\xi(x_c) \in [e_{x_c} + \bar{\gamma}^-(x_c), c_{x_c} + \bar{\gamma}^+(x_c)]$.

We sometimes use interval bounded schedules for requirement events as well, where $\xi(x_r) = t = [t, t]$ for some requirement event x_r assigned to time t.

We once again revise our definition of situations, Ω_v, to reflect the impact of the variable-delay function on the space of observations.

Definition 18. Variable-Delay Situations

For an STNU S with k contingent constraints $\langle e_1, c_1, l_1, u_1 \rangle, \cdots, \langle e_k, c_k, l_k, u_k \rangle$ and variable-delay function $\bar{\gamma}$, each *variable-delay situation*, ω_v, represents a possible set of *observed* values for all links in S, $\omega = (\omega_{v1}, \cdots, \omega_{vk})$. The {space of situations} for S, Ω_v, is $\Omega_v = [e_1, c_1] + [\bar{\gamma}_1^-, \bar{\gamma}_1^+] \times \cdots \times [e_k, c_k] + [\bar{\gamma}_k^-, \bar{\gamma}_k^+]$.

We see that the space of observations has likewise grown in the transition to variable observation delay. If $\bar{\gamma}^- < \bar{\gamma}^+$, Ω_v for variable observation delay is strictly larger than Ω_f for fixed-observation delay and Ω for vanilla STNUs.

Like the fixed-delay function for fixed-delay controllability, the variable-delay function relates an observation delay to a contingent event, independent of other events. We take a similar approach to defining variable-delay controllability, relative to fixed-delay controllability.

Definition 19. Variable-Delay Controllability

An STNU S is *variable-delay controllable* with respect to a variable-delay function, $\bar{\gamma}$, if and only if for the space of situations, Ω_v, there is an $\mathcal{S}$ that produces a satisfying schedule for requirement events during execution, ξ.

Determining whether a given variable-delay STNU, S, is variable-delay controllable has two components [1]. The first is to derive a fixed-delay STNU, S', with fixed-observation delay, γ, that is equivalent with respect to controllability. The second is to show that S' is fixed-delay controllable. Below, we reiterate the claims of [1], demonstrating how to derive S' from S that is equivalent with respect to controllability. We first demonstrate how to transform the contingent links from S to S', and demonstrate their correctness with respect to observation spaces, before following up with transformations to the requirement links to maintain the same scheduling semantics in S'.

For the following lemmas, let x_c be a contingent link in S, where $x_c \in [l, u]$ and variable-delay function $\bar{\gamma}(x_c)$. Let x'_c be the transformed contingent link in S' with fixed-delay function, $\gamma(x'_c)$.

Note that we receive $\psi(x_c)$ from Nature, but make the assignment $\xi(x'_c)$ in the dispatchable form of S'. To be clear, while $\xi(x_c)$ is an interval, $(\mathbb{R} \cup \infty) \times (\mathbb{R} \cup \infty)$, $\xi(x'_c)$ is in $\mathbb{R}$. For a fixed interval, e.g. $\psi(x_c) \in [t, t]$, we sometimes employ an equivalent representation, $\xi(x_c) = t$.

Additionally, we sometimes apply $-$ and $+$ superscripts to l and u to denote the earliest and latest times respectively that an assignment at those bounds could be observed. For instance, the relationship in Definition 16 simplifies to,

$$\psi(x_c) = [l + \bar{\gamma}^-(x_c), u + \bar{\gamma}^+(x_c)] \tag{4.1}$$

$$\psi(x_c) = [l^-(x_c), u^+(x_c)] \tag{4.2}$$

Figure 4-1: We visualize the relationship between realized assignments across S and S'. In this example, each horizontal line is a timeline monotonically increasing from left to right. Dashed lines represent observation delays. We see how an assignment in S, $\xi(\mathrm{x_c})\$$, realized observation delay, $g(x_c)$, and an observation in S, $\psi(x_c)$, lead to an assignment in S', $\xi(\mathrm{x'_c})\$$.

Lastly, we need a means to compare observation spaces if we are to transform variable-delay to fixed-delay STNUs.

Definition 20. Observation Space Mapping

Let μ be a mapping from an assignment to a situation, $\mu : \xi \to \omega$. To say that $\mu(x'_c) \subseteq \omega_v(x_c)$ means that, for any assignment of x'_c in S', there is an equivalent situation in S for x_c.

For the transitions below, it is a *valid observation space mapping*, if we can show that $\mu(x'_c) \subseteq \omega_v(x_c)$. If so, it is guaranteed that any assignment in the observation space of x'_c also has a valid assignment in the observation space of x_c.

We now have the necessary vocabulary and notation to step through the transformations from S to S'. These lemmas were first presented in [1].

Definition 21. Variable-Delay to Fixed-Delay Transformations

The *variable-delay to fixed-delay transformations* define a set of observation space mappings, where there are valid observation space mappings for all the contingent constraints in S' to S.

Thus, if there is a satisfying $\mathcal{S}$ for the fixed-delay observation space of S', it is guaranteed to simultaneously satisfy any situation in the variable-delay observation space, Ω_v, of S.

Lemma 1. *For any contingent event $x_c \in X_c$ in S, if $\bar{\gamma}^-(x_c) = \bar{\gamma}^+(x_c)$, we emulate $\bar{\gamma}(x_c)$ in S' using $\gamma(x'_c) = \bar{\gamma}^+(x_c)$.*

Proof. We translate an already fixed-bounded observation delay in the form of $\bar{\gamma}(x_c)$ to the equivalent fixed-delay function, $\gamma(x'_c)$, thus $\omega_f(x'_c) = \omega_v(x_c)$. □

Lemma 2. *For any contingent event $x_c \in X_c$, $\bar{\gamma}^+(x_c) = \infty$, we emulate $\bar{\gamma}(x_c)$ in S' as $\gamma(x'_c) = \infty$.*

Proof. There are projections where we would not receive information about x_c, therefore we have to act as if we *never* receive an observation of x_c. Any $\mathcal{S}$ that works when we do not receive information about x_c would also work when do receive an observation if we choose to ignore the observation.

None of our decisions depend on $\xi(x'_c)$, thus no observation space mapping to S is necessary. □

Lemma 3. *If $u - l \leq \bar{\gamma}^+(x_c) - \bar{\gamma}^-(x_c)$, we emulate $\bar{\gamma}(x_c)$ in S' using $\gamma(x'_c) = \infty$.*

Proof. We can ignore observations of x_c because they are not guaranteed to narrow where $\xi(x_c)$ was assigned in the range $[l, u]$.

Let α be the range of $\psi(x_c)$ when $\xi(x_c) \in [l, l]$. Let β be the range of $\psi(x_c)$ when $\xi(x_c) \in [u, u]$. By Equation 4.1,

$$\alpha = [l^-(x_c), l^+(x_c)]$$
$$\beta = [u^-(x_c), u^+(x_c)]$$

We can show that $u^-(x_c) \leq l^+(x_c)$.

$$u - l \leq \bar{\gamma}^+(x_c) - \bar{\gamma}^-(x_c)$$
$$u + \bar{\gamma}^-(x_c) \leq l + \bar{\gamma}^+(x_c)$$
$$u^-(x_c) \leq l^+(x_c)$$

The lower bound of β is less than the upper bound of α, thus $\alpha \cap \beta$. An observation $\psi(x_c) \in [u^-(x_c), l^+(x_c)]$ could be the result of $\xi(x_c) = [l, l]$, $\xi(x_c) = [u, u]$, or any value $\xi(x_c) \in [l, u]$. Observations provide no information about the underlying contingent constraint, therefore we ignore $\psi(x_c)$.

None of our decisions depend on $\xi(x'_c)$, thus no observation space mapping to S is necessary. □

Lemma 4. *If $u - l \geq \bar{\gamma}^+(x_c) - \bar{\gamma}^-(x_c)$, we can emulate $\bar{\gamma}(x_c)$ under minimal information by replacing the bounds of x_c with $x'_c \in [l^+(x_c), u^-(x_c)]$ and letting $\gamma(x'_c) = 0$.*

Proof. Under Lemma 4, observations $\psi(x_c)$ are guaranteed to narrow the range of $\xi(x_c)\$.

We have the same ranges for α and β as in Lemma 3, however we can show that $u^-(x_c) \geq l^+(x_c)$ instead.

$$u - l \geq \bar{\gamma}^+(x_c) - \bar{\gamma}^-(x_c)$$

$$u + \bar{\gamma}^-(x_c) \geq l + \bar{\gamma}^+(x_c)$$

$$u^-(x_c) \geq l^+(x_c)$$

Thus, receiving an observation is guaranteed to narrow the derived range of $\xi(x_c)$. The transformation tightens the range of x_c' to one where there is maximum ambiguity of the assignment of x_c while guaranteeing an execution strategy for any assignment of $x_c \in [l, u]$.

Based on the derivations above, it is clear that $\mu(x_c')$ maps to the observation space where there is ambiguity as to the projection of $\xi(x_c) \in [l, u]$. We must also show that $\mu(x_c')$ has mappings to the extrema of $\xi(x_c)$. We start with the earliest $\xi(x_c')$.

$$\xi(x_c') = l^+(x_c) = l + \bar{\gamma}^+(x_c)$$

We show that that this assignment of $\xi(x_c')$ can be modeled as the following observation in S.

$$\psi(x_c) \in [l + \bar{\gamma}^-(x_c), l + \bar{\gamma}^+(x_c)]$$

$$\psi(x_c) \in [l, l] + \Gamma(x_c)$$

It is possible that $\xi(x_c) = [l, l]$. As such, all observations in $\psi(x_c)$ may share the same execution strategy because the underlying temporal constraints depend on $\xi(x_c)$, not $\psi(x_c')$ or $\Gamma(x_c)$. We may expand the range of the observation space when we map to S with $\mu(x_c')$.

$$\mu : l^+(x_c) \to \omega_v(x_c)$$

$$\omega_v(x_c) = [l + \bar{\gamma}^-(x_c), l + \bar{\gamma}^+(x_c)]$$

We see that μ has a valid observation space mapping to the minimum of the range of $\omega_v(x_c)$. We use the same argument for the maximum.

$$\xi(x_c') = u + \bar{\gamma}^-(x_c)$$

Observations anywhere in $[u + \bar{\gamma}^-(x_c), u + \bar{\gamma}^+(x_c)]$ may share execution strategies because, it is possible that in all cases, $\xi(x_c) = [u, u]$. We may then expand the range of the observation space when we map to S.

$$\mu : u^-(x_c) \rightarrow \omega_v(x_c)$$

$$\omega_v(x_c) = [u + \bar{\gamma}^-(x_c), u + \bar{\gamma}^+(x_c)]$$

Thus, $\mu(x_c')$ maps to the maximum of the range of $\omega_v(x_c)$. The transition creates assignments in S' that map to the entire $\omega_v(x_c)$ in S.

$\square$

This concludes the modifications required to transform a contingent event $x_c \in X_c$ in S to its equivalent $x_c' \in X_c$ in S'. What remains is to address the transformation of requirement links, $x_r \in X_r$, in S such that their transformed equivalents, $x_r' \in X_r$ in S', express the same execution semantics in S' as they did in S. We will demonstrate the correctness of the transformations after Lemma 6.

Lemma 5. *If we have contingent link $X \Rightarrow C$ with duration $[l, u]$, outgoing requirement link $C \rightarrow Z$ with duration $[u, v]$ with an unobservable C, and contingent link $C \Rightarrow Y$ with range $[\bar{\gamma}^-(x_c), \bar{\gamma}^+(x_c)]$, we can emulate the role of the original requirement link during execution with a new link $Y \rightarrow Z$ with bounds $[u - max(\bar{\gamma}^-(x_c), XY - u), v - min(\bar{\gamma}^+(x_c), XY - l)]$, where XY is the true duration of $X \Rightarrow Y$.*

Proof. See Figure **??**c for reference. From an execution perspective, X and Y are the only events that can give us any information that we can use to reason about when to execute Z (since C is wholly unobservable).

If we execute Z based on what we learn from Y, then we use our information from Y to make inferences about the true durations of $X \Rightarrow C$ and $C \Rightarrow Y$ based on $X \Rightarrow Y$. We know that the lower-bound of $C \Rightarrow Y$ is at least $XY - b$ and that its upper-bound is at most $XY - a$. But we also have the a priori bounds on the contingent link that limit its range to $[\bar{\gamma}^-, \bar{\gamma}^+]$. Taken together, during execution we can infer that the true bounds of $C \Rightarrow Y$ are $[max(\bar{\gamma}^-, XY - b), min(\bar{\gamma}^+, XY - a)]$. Since we have bounds only on Z's execution in relation to C, we can then infer a requirement link $Y \rightarrow Z$ with bounds $[u - max(\bar{\gamma}^-, XY - b), v - min(\bar{\gamma}^-, XY - a)]$.

If we try to execute Z based on information we have about X, we must be robust to any possible value assigned to $X \Rightarrow C$. This means that we would be forced to draw a requirement link $X \rightarrow Z$ with bounds $[u + b, v + a]$. But we know that $u - max(\bar{\gamma}^-, XY - b) \leq u + b - XY$ and $v - min(\bar{\gamma}^-, XY - a) \geq v + a - XY$, which means that the bounds we derived from Y are at least as expressive as the bounds that we would derive from X. $\square$

Since we have a local execution strategy that depends on the real value of XY, we can try to apply this strategy to the contingent link that we restricted in Lemma 4, in order to repair the

remaining requirement links.

Lemma 6. *If we have an outgoing requirement link $C \to Z$ with duration $[u, v]$, where C is a contingent event, we can emulate the role of the original requirement link by replacing its bounds with $[u - \bar{\gamma}^-(x_c), v - \bar{\gamma}^+(x_c)]$.*

Proof. See Figure **??**d for reference. If we directly apply the transformation from Lemma 5 and Figure **??**c to our original STNU, we introduce complexity through the need to reason over *min* and *max* operations in our link bounds. However, from Lemma 4, we know that in a controllability evaluation context, it is acceptable for us to simplify the $X \Rightarrow Y$ link to a stricter range of $[a + \bar{\gamma}^+, b + \bar{\gamma}^-]$, instead of $[a + \bar{\gamma}^-, b + \bar{\gamma}^+]$. This means that for the purpose of evaluating controllability, we can assume $a + \bar{\gamma}^+ \leq XY \leq b + \bar{\gamma}^-$. When we evaluate the requirement link $Y \to Z$, we see $max(\bar{\gamma}^-, XY - b) = \bar{\gamma}^-$ and $min(\bar{\gamma}^+, XY - a) = \bar{\gamma}^+$. This gives us bounds of $[u - \bar{\gamma}^-, v - \bar{\gamma}^+]$ for the $Y \to Z$ requirement link as seen in Figure **??**d. $\square$

Lemma 6 handles outgoing requirement edges connected to contingent events. In addition, we must handle incoming edges.

Corollary 6.1. *If we have an incoming requirement link $Z \to C$ with duration $[u, v]$, where C is a contingent event, we can replace the bounds of the original requirement link with $[u + \bar{\gamma}^+(x_c), v + \bar{\gamma}^-(x_c)]$.*

Proof. A requirement link $Z \to C$ with bounds $[u, v]$ can be immediately rewritten as its reverse $C \to Z$ with bounds $[-v, -u]$. After reversing the edge, we can apply Lemma 6 to get $Y \to Z$ with bounds $[-v - \bar{\gamma}^-, -u - \bar{\gamma}^+]$, which we can reverse again to get $Z \to Y$ with bounds $[u + \bar{\gamma}^+, v + \bar{\gamma}^-]$. $\square$

After applying Lemma 4, despite the limited expected range of assignments in x_c' in S' compared to x_c in S, we can show that Lemma 6 guarantees a satisfying schedule for any $\psi(x_c) \in [l^-(x_c), u^+(x_c)]$ using an S that employs *buffering* and *imagining* contingent events.

Definition 22. Buffering

Buffering a contingent event x_c is an execution strategy where, if x_c is observed earlier than the lower bound of the observation space $\psi(x_c) < \omega_f^-(x_c')$, we assign $\xi(x_c')$ to the lower bound of the observation space, $\xi(x_c') = \omega_f^-(x_c')$.

Definition 23. Imagining

Imagining a contingent event x_c is an execution strategy where, if x_c is observed later than the upper bound of the observation space, $\psi(x_c) > \omega_f^+(x_c')$, we assign $\xi(x_c')$ to the upper bound of the observation space, $\xi(x_c') = \omega_f^+(x_c')$.

Lemma 7. *If S' is fixed-delay controllable after applying Lemmas 4, 5, and 6 to contingent event Y with following requirement event Z, there is a valid S for any observation in the observation space of S, $\omega_v(Y) = [a^-(Y), b^+(Y)]$.*

Proof. We first note the observation space of S' is a subinterval of the original observation space of S, $\omega_f(Y') \subset \omega_v(Y)$, and there are two distinct ranges of observations that are not in $\omega_f(Y')$.

$$\omega_f(Y') = [a + \bar{\gamma}^+(Y), b + \bar{\gamma}^-(Y)]; \ \omega_v(Y) = [a + \bar{\gamma}^-(Y), b + \bar{\gamma}^+(Y)]$$

$$\omega_f(Y') \not\supset [a + \bar{\gamma}^-(Y), a + \bar{\gamma}^+(Y)) \ (\text{"Early" observations})$$

$$\omega_f(Y') \not\supset (b + \bar{\gamma}^+(Y), b + \bar{\gamma}^+(Y)] \ (\text{"Late" observations})$$

We address the early observations first. The range of early assignments of $\xi(Y)$ in S that we care about are the ones that could produce an observation $\psi(Y) \leq a + \bar{\gamma}^+(Y)$, which is $\xi(Y) = [a, a + (\bar{\gamma}^+(Y) - \bar{\gamma}^-(Y))]$. We rewrite the range of early assignments as $\xi(Y) = a + (\bar{\gamma}^+(Y) - \bar{\gamma}^-(Y)) - \epsilon$, where $0 \leq \epsilon \leq (\bar{\gamma}^+(Y) - \bar{\gamma}^-(Y))$. By the semantics of S, the range of assignments of $\xi(Z)$ is then,

$$\xi(Z) = [a + (\bar{\gamma}^+(Y) - \bar{\gamma}^-(Y)) - \epsilon, a + (\bar{\gamma}^+(Y) - \bar{\gamma}^-(Y)) - \epsilon] + [u, v]$$

$$\xi(Z) = [a + u + (\bar{\gamma}^+(Y) - \bar{\gamma}^-(Y)) - \epsilon, a + v + (\bar{\gamma}^+(Y) - \bar{\gamma}^-(Y)) - \epsilon]$$

The earliest assignment of Y' in S' is $\xi(Y') = a + \bar{\gamma}^+(Y)$. By the semantics of S', the range of assignments of $\xi(Z')$ is then,

$$\xi(Z') = [a + \bar{\gamma}^+(Y), a + \bar{\gamma}^+(Y)] + [u - \bar{\gamma}^-(Y), v - \bar{\gamma}^+(Y)]$$

$$\xi(Z') = [a + u + (\bar{\gamma}^+(Y) - \bar{\gamma}^-(Y)), a + v]$$

We see that $\xi(Z') \subseteq \xi(Z)$ for any ϵ, meaning the execution strategy when $\xi(Y') = a + \bar{\gamma}^+(Y)$ results in a valid assignment of $\xi(Z)$ for all early observations of $\xi(Y)$. We are safe to buffer early observations to $\xi(Y') = a + \bar{\gamma}^+(Y)$.

We use the same argument for imagining late observations. The range of late assignments of $\xi(Y)$ in S that we care about are the ones that could produce an observation $\psi(Y) \geq b + \bar{\gamma}^-(Y)$, which is $\xi(Y) = b - (\bar{\gamma}^+(Y) - \bar{\gamma}^-(Y)) + \epsilon$. By the semantics of S, the range of assignments of $\xi(Z)$ is then,

$$\xi(Z) = [b - (\bar{\gamma}^+(Y) - \bar{\gamma}^-(Y)) + \epsilon, b - (\bar{\gamma}^+(Y) - \bar{\gamma}^-(Y)) + \epsilon] + [u, v]$$

$$\xi(Z) = [b + u - (\bar{\gamma}^+(Y) - \bar{\gamma}^-(Y)) + \epsilon, b + v - (\bar{\gamma}^+(Y) - \bar{\gamma}^-(Y)) + \epsilon]$$

The last assignment of Y' in S' is $\xi(Y') = b + \bar{\gamma}^-(Y)$. By the semantics of S', the range of

assignments of $\xi(Z')$ is then,

$$\xi(Z') = [b + \bar{\gamma}^-(Y), b + \bar{\gamma}^+(Y)] + [u - \bar{\gamma}^-(Y), v - \bar{\gamma}^+(Y)]$$
$$\xi(Z') = [b + u, b + v - (\bar{\gamma}^+(Y) - \bar{\gamma}^-(Y))]$$

We see that $\xi(Z') \subseteq \xi(Z)$ for any ϵ, meaning the execution strategy when $\xi(Y') = b + \bar{\gamma}^-(Y)$ results in a valid assignment of $\xi(Z)$ for all late observations of $\xi(Y)$. In practice, there is no reason to wait until after $\psi(Y) = b + \bar{\gamma}^-(Y)$ to receive a late observation. As soon as we see the clock has reached $b + \bar{\gamma}^-(Y)$, we are safe to imagine that $\psi(Y)$ has been received. $\qquad\square$

We can examine a concrete example of Lemmas 4, 5, and 6 to show equivalence in the transformation from Figure **??**a to **??**d. We start by building an example of **??**a. Let $X \xRightarrow{[2,5]} C$ with $\bar{\gamma}(C) \in [1,2]$ and $C \xrightarrow{[11,20]} Z$. If we learn of event C at time 4, then one possibility is that the realized duration of C could have been 2 with an observation delay of 2. In this case, event Z must be executed in $[13, 22]$. However, if the realized duration of C were 3 with an observation delay of 1, then Z would fall in $[14, 23]$. Given we cannot distinguish between the possibilities, we take the intersection of the intervals, yielding $Z \in [14, 22]$. Likewise, if we learn of C at time 6, then C could have been realized at time 5 with an observation delay of 1 or it could have been realized at time 4 with an observation delay of 2. In the first case, Z must then fall in $[16, 25]$, while in the second, Z would fall in $[15, 24]$. The intersection yields $[16, 24]$.

By the semantics represented in Figure **??**d, we can build an equivalent network with $\gamma(Y) = 0$ by setting $X \xRightarrow{[4,6]} Y$ and $Y \xrightarrow{[10,18]} Z$. If Y is observed at time 4, Z must be executed in $[14, 22]$. If Y is observed at time 6, Z then must be executed in $[16, 24]$. The execution semantics for both cases match the equivalent networks from **??**a described above.

4.4 Dynamic Scheduling through Real-Time Execution Decisions

An STNU, S, that exhibits dynamic controllability can be *scheduled* dynamically (or *online*). At a high-level, dynamic scheduling is the process of mapping the history of event assignments to the execution time of future free events. We follow the scheduling work by Hunsberger [34], [37], which describes an $O(N^3)$ procedure, FAST-EX, for dynamic scheduling of STNUs. At its core is the notion of *Real-Time Execution Decisions* (RTEDs), which map a timepoint to a set of requirement events to be executed and are generated based on *partial schedules* of STNUs being executed. WAIT decisions may also be produced, reflecting the need to wait for the assignment of a contingent

event before continuing. RTED-based scheduling applies a dynamic programming paradigm by first creating a dispatchable form of temporal constraints offline, updating the dispatchable form as the partial schedule is updated online, and querying the dispatchable form online to quickly find the next free event to schedule [38].

Definition 24. Real-Time Execution Decisions [38]

A *Real-Time Execution Decision* is a two-tuple $\langle t, \chi \rangle$, where:

- t is a time with domain $\mathbb{R}$,

- χ is a set of $x_r \in X_r$ to be executed at time t

Definition 25. Partial Schedule

A *Partial Schedule* for an STNU is a mapping $\xi : t \rightarrow \mathbb{R}$, where t is a proper subset of timepoints in X with domain $\mathbb{R}$.

A partial schedule is the set of assignments *so far* during execution. They can be represented by $\xi_t = \{(A, \xi(A)\}$ notation, meaning that by some time t, event A has been scheduled at time $\xi(A)$. By definition, all STNU executions begin with a partial schedule $\xi_0 = \emptyset$, meaning no timepoints have been assigned at time 0. As timepoints are executed, the partial schedule grows. For example, if event A is assigned to 1 and B to 2, then at time \$2, $\xi_2 = \{(A, 1), (B, 2)\}$.

The dispatchable form employed by FAST-EX is the *AllMax* distance graph, which is produced by the Morris $O(N^4)$ DC-checking procedure [27].

Definition 26. AllMax Distance Graph [26]

The *AllMax* distance graph is a distance graph exclusively consisting of unlabeled and upper-case edges.

The key idea of FAST-EX is maintaining accurate distances from a zero point, Z, of the graph to all events. At the outset of execution, all events from S are present as nodes in *AllMax*. As contingent events are observed, *AllMax* performs update steps using Dijkstra Single Source/Sink Shortest Path (SSSP) to maintain distances to unexecuted events, while also removing executed events. We include pseudo-code of the update step in Figure 1.

Given an accurate distance matrix, we find the next RTED as follows. Let U_x be the set of unexecuted free timepoints. If U_x is empty, then the RTED is to `WAIT`. Otherwise, we find the lower bound of the earliest executable time point and the set of executable events associated with it.

$$t = \min\{-D(X, Z) \mid X \in U_x\}$$

$$\chi = \{X \in U_x \mid -D(X, Z) = t\}$$

Input: Time t; Set of newly executed events $\texttt{Exec} \subseteq X_e \cup X_r$; AllMax Graph G; Distance matrix D, where $D(A, B)$ is the distance from A to B

Output: Updated D

   FAST-EX Update:

      for *each continent event $C \in \texttt{Exec}$* **do**

         Remove each upper-case edge, $Y \xrightarrow{C:-w} A$, labled by C;

         Replace each edge from Y to Z with the strongest replacement edge;

      end

      for *each event $E \in \texttt{Exec}$* **do**

         Add lower-bound edge $E \xrightarrow{-t} Z$;

      end

      For each event X, update $D(X, Z)$ using Dijkstra Single-Sink Shortest Paths;

      for *each event $E \in \texttt{Exec}$* **do**

         Add upper-bound edge $Z \xrightarrow{t} E$;

      end

      For each event X, update $D(Z, X)$ using Dijkstra Single-Source Shortest Paths;

  Algorithm 1: Algorithm for updating distances for all events in relation to Z upon the execution of an event. Adapated from [3], Fig. 19.

We cannot execute events in the past. Let $\texttt{now}$ be the current time, i.e. the last timepoint captured in ξ. It is possible that $t \leq \texttt{now}$, in which case we must reassign t to guarantee that $t > \texttt{now}$. To do so, we update t as follows, where t_U is earliest *upper* bound of the executable timepoints,

$$t_U = \min\{D(Z, X) \mid X \in U_x\}$$

$$t = \frac{\texttt{now} + t_U}{2}$$

So long as $t_U > \texttt{now}$, we know that the reassignment of t ensures $t > \texttt{now}$.

THIS PAGE INTENTIONALLY LEFT BLANK

Chapter 5

# Dynamic Scheduling with Delayed Event Monitoring

Figure 5-1: From a variable-delay STNU to scheduling decision

This section represents the start of contributions from this paper. Our aim is to describe *delay scheduling*, an RTED-based strategy for dispatching free events in the case where there is set-bounded observation delay for contingent events. As will be shown, delay scheduling is an extension to dynamic scheduling using the execution strategy as outlined in the variable-delay controllability checking procedure.

We must address three gaps: first, we reconcile

At the conclusion of this Section, we extend our approach to scheduling variable-delay STNUs by introducing an optional procedure that addresses a shortcoming in the semantics of scheduling a variable-delay STNU. The shortcoming takes the form of potentially unnecessary wait times that are added after receiving contingent event assignments, extending the makespan of procedures. We present a generate-and-test algorithm to partially mitigate said shortcomings.

Bhargava et al. [1] addressed this ambiguity in contingent event assignment by first transforming the VDC STNU into a controllability-equivalent fixed-delay STNU. With fixed observation delay, we *do* have the guarantee that we learn the exact assignment of contingent events (so long as the

observation delay is not infinite). Thus, scheduling a fixed-delay STNU only differs from scheduling a vanilla STNU in that we must subtract a fixed observation delay when we make contingent event assignments. Otherwise, the dispatchable form is the same as in the case of a vanilla STNU, and we can choose any STNU scheduling algorithm to generate execution decisions.

The flow from variable-delay STNU to fixed-delay STNU to dispatchable form may appear sufficient to enable scheduling of variable-delay STNUs, but we must contend with a novel issue: the execution spaces of the original variable-delay STNU and its transformed fixed-delay equivalent are mismatched. Nature is obliged to respect the uncertainties of the original variable-delay STNU. As will be shown later, the fixed-delay equivalent reduces the execution space to make the controllability check tractable. As such, we may receive observations outside the range of the contingent links in the fixed-delay STNU, which we must reconcile with the dispatchable form. See Figure 5-1 for an overview of the information flow in scheduling a variable-delay STNU.

5.1 Scheduling with Variable-Observation Delay

To solidify the process of scheduling a variable-delay STNU, consider the following analogy.

> Alex wants to go hiking in the woods. The area is unfamiliar to them, so they ask their friend, Sam, who hiked these trails a long time ago, to give them directions to traverse from the trailhead to a particularly spectacular overlook. Sam has a working idea of the trail map, but their memory is imperfect. Regardless, they guarantee Alex that their directions will lead Alex to the overlook even if the woods have changed over the years. Sam writes down directions like "turn left after 500 meters at the giant oak tree" and "turn right after 100 meters when you see the brook." Alex knows that Nature will not necessarily obey Sam's directions. They may observe a giant oak tree earlier than expected, so they must then wait to take the next trail going left. Or the brook may have dried up, so they imagine they saw one near where Sam thought it would be and take the next right. While hiking, Alex is charged with reconciling Sam's directions with their own observations. Even though they may identify the landmarks in Sam's directions earlier or later than expected, their actions will need to follow Sam's instructions to maintain the guarantee of reaching the overlook.

In our analogy, S models the current state of the hiking trails and the full range of projections, while S' is Sam's working memory of them. Sam's directions are the execution strategy described by the AllMax graph we get by checking the fixed-delay controllability of S'. Observations of Nature obey S. Alex is charged with reconciling their observations from S with Sam's hiking directions from S'. The analogy ends here, though, as the math and logic of temporal reasoning do not

neatly translate into hiking. Luckily, we have more information than Alex. Unlike human memory, which is untrustworthy and irrational, the fixed-delay STNU, S', is created by a set of Lemmas with deterministic outcomes. As such, we have the means to interpret how observations in S *would appear* in S', which will be critical in adapting our fixed-delay execution strategy in response to variable observation delay.

Our key challenge for scheduling an STNU with variable observation delay is reconciling observations from S with the dispatchable form from S'.

5.2 Recording Contingent Event Assignments

During execution, we observe the outcome of contingent events $\psi(x_c)$ in S, but we make assignments in the dispatchable form of $\xi(x'_c)$ in S'. Despite being equivalent with respect to controllability, the bounds of contingent links x_c in S and x'_c in S' are not equivalent.

We need a modified procedure for contingent event assignments that wraps FAST-EX. No modifications to FAST-EX are necessary to schedule fixed-delay STNUs because checking FDC includes the procedure of creating the same AllMax graph that FAST-EX requires.

We now present our strategy for recording observations during execution as derived from the transformations outlined in Section 4.3.

Lemma 8. *For any contingent event, $x_c \in S$ or $x'_c \in S'$, observing x_c at time $t \in [l^-(x_c), u^+(x_c)]$ fixes the observation to $\psi(x_c) = [t, t] = t$.*

Proof. Prior to execution, observations are defined as set-bounded intervals from the earliest possible observation at $l^-(x_c)$ to the last possible observation at $u^+(x_c)$. Receiving an observation $\psi(x_c) = t$ during execution eliminates all members of the pre-execution interval except t. $\square$

Lemma 9. *For any contingent event $x'_c \in X_c$ in fixed-delay controllable S', if $\gamma(x'_c) = \infty$, we mark the event executed but do not assign $\xi(x'_c)$ in the dispatchable form of S'.*

Proof. If we are scheduling a fixed-delay STNU, S', that is already known to be fixed-delay controllable, an execution strategy must exist that is independent of the assignment of $\xi(x'_c)$ when $\gamma(x'_c) = 0$. We are not required to record $\xi(x'_c)$ when $\gamma(x'_c) = \infty$ to guarantee controllability and may safely ignore it.

We mark the event executed to prevent it from appearing in future RTEDs.

$\square$

Lemma 9 may be applicable to any contingent events, $x'_c \in X_c$ in S' that were transformed from the variable-delay form S using Lemmas 1, 2, or 3.

Lemma 10. *For any contingent event $x'_c \in X_c$ in fixed-delay controllable S', if $\gamma(x'_c) \in \mathbb{R}$, we assign $\xi(x'_c) = \psi(x_c) - \gamma(x'_c)$ in the dispatchable form of S'.*

Proof. The central challenge of checking fixed-delay controllability is determining that an execution strategy exists that allows an agent to wait an additional $\gamma(x'_c)$ time units after a contingent event has been assigned to learn its outcome. Importantly, the γ function is not used to modify the edges of the labeled distance graph, which are derived from the constraints $r \in R_e \cup R_c$ in S'.

As $\gamma(x'_c)$ resolves to a known and finite value, we can derive the true value of $\xi(\mathrm{x'}_c)\$$ to be assigned in the labeled distance graph. Contingent event assignments, $\xi(\mathrm{x'}_c)\$$, are recorded in the labeled distance graph as follows, where $\psi(x_c)$ is the resolved observation,

$$\xi(x'_c) = \psi(x_c) - \gamma(x'_c) \tag{5.1}$$

$\square$

Next, in comparing the bounds of x_c and x'_c when $u - l \geq \bar{\gamma}^+(x_c) - \bar{\gamma}^-(x_c)$, $x'_c \in [l^+(x_c), u^-(x_c)]$ (Lemma 4) there are three regimes of observations of $\psi(x_c)$ we must consider:

1. $\psi(x_c) \in [l^-(x_c), l^+(x_c))$, ie. strictly earlier than the range of $\xi(x'_c)$,

2. $\psi(x_c) \in [l^+(x_c), u^-(x_c)]$, ie. the range equivalent to x'_c, and

3. $\psi(x_c) \in (u^-(x_c), u^+(x_c)]$, ie. strictly later than the range of $\xi(x'_c)$.

Nature decides in which regime we receive $\psi(x_c)$. We are faced with the unique challenge of deciding how to act when Nature selects an $\psi(x_c)$ that fails to follow the constraints of S', eg. $\psi(x_c) < l^+(x_c) \vee \psi(x_c) > u^-(x_c)$, which would lead to an assignment, $\xi(x'_c)$, in the first or third regimes above. In plainer words, the contingent links of S and S' do not have the same constraints. We make assignments in S', but we receive observations from S. We need to decide how to act when we observe a contingent event earlier or later than we expect according to S', because if we blindly assigned $\xi(x'_c)$ outside its constraints from S', we lose the guarantee of controllability. Our only choice is to find a strategy to assign x'_c that respects the constraints of S', despite observing x_c earlier or later than expected. We do so by reasoning over the possible *range* of assignments, $\xi(x_c)$, that could have led to a particular observation, $\psi(x_c)$. What we find is that, due to the uncertainty in observation delay, we are allowed to *modify* our assignment of $\xi(x'_c)$ to ensure it respects S'. We present two modification strategies for addressing the first and third cases, which we call *buffering* and *imagining* respectively.

We first address the case where $\psi(x_c) < l^+(x_c)$.

Lemma 11. *If a contingent event, $x_c \in X_c$, is observed earlier than the bounds of x'_c in S' for a fixed-delay controllable S', $\psi(x_c) < l^+(x_c)$, we perform a buffering operation by letting $\xi(x'_c) = l^+(x_c)$ in S'.*


Figure 5-2: Here, we show how the combination of $\xi(x_c)$ and $\bar{\gamma}(x_c)$ lead to an assignment of $\xi(x'_c)$ in S'. We see the range $\alpha \in [l, l + \bar{\gamma}^+(x_c) - \bar{\gamma}^-(x_c)$ representing the earliest and latest assignments of $\xi(x_c)\$$ that could result in $\psi(x_c) \in \xi(x'_c) \in [l^+(x_c), 1\hat{}+(x_c)]\$$. The grey region represents the range of possible observation delays, $\bar{\gamma}(x_c)$, supporting $\xi(x'_c) \in [l^+(x_c), l^+(x_c)]$.

Proof. To demonstrate why buffering is sound, we compare the bounds of x_c in S and x'_c in S' to show that our execution strategy for $\xi(x'_c)$ is applicable to any $\xi(x_c) \in [l, l^+(x_c)]$.

We know that S' is fixed-delay controllable when $\xi(x'_c) \in [l^+(x_c), u^-(x_c)]$. Consider an observation at the lower bound of $\$\xi(x'_c)$, $\psi(x_c) = l^+(x_c)$. We can discern the range of possible assignments of x_c in S (Using Lemma **??** to rewrite $o(x_c) = l^+(x_c)$ as $o(x_c) = [l^+(x_c), l^+(x_c)]$).

$$\psi(x_c) = \xi(x_c) + \bar{\gamma}(x_c)$$
$$\xi(x_c) = \psi(x_c) - \bar{\gamma}(x_c)$$
$$\xi(x_c) = [l^+(x_c), l^+(x_c)] - [\bar{\gamma}^-(x_c), \bar{\gamma}^+(x_c)]$$
$$\xi(x_c) = [l, l + (\bar{\gamma}^+(x_c) - \bar{\gamma}^-(x_c))]$$

Let $\alpha = [l, l + (\bar{\gamma}^+(x_c) - \bar{\gamma}^-(x_c))]$ for this Lemma.

Given S' is fixed-delay controllable, there must exist an execution strategy when $\xi(x'_c) = l^+(x_c)$, which entails the same execution strategy applies for any assignment of $\xi(x_c) \in \alpha$. Thus, during

execution, if we can show that $\xi(x_c) \subseteq \alpha$, we can safely act as if $\xi(x_c') = l^+(x_c)$.

Now, let $\psi(x_c) = l^+(x_c) - \epsilon$ for some small, positive ϵ. Accordingly, it is the case that $\xi(x_c)$ must fall in the range,

$$\xi(x_c) = [(l^+(x_c) - \epsilon) - [\bar{\gamma}^-(x_c), \bar{\gamma}^+(x_c)]$$

$$\xi(x_c) = [l^+(x_c) - \epsilon, l^+(x_c) - \epsilon] - [\bar{\gamma}^-(x_c), \bar{\gamma}^+(x_c)]$$

$$\xi(x_c) = [l - \epsilon, l + (\bar{\gamma}^+(x_c) - \bar{\gamma}^-(x_c)) - \epsilon]$$

Of course, $\xi(x_c)$ must respect the original bounds of x_c, $x_c \in [l, u]$.

$$\xi(x_c) = [l - \epsilon, l + \bar{\gamma}^+(x_c) - \bar{\gamma}^-(x_c) - \epsilon] \cap [l, u]\xi(x_c) \qquad = [l, l + (\bar{\gamma}^+(x_c) - \bar{\gamma}^-(x_c)) - \epsilon]$$

Let $\beta = [l, l + (\bar{\gamma}^+(x_c) - \bar{\gamma}^-(x_c)) - \epsilon]$ for this Lemma. See Figure 5-2 for a visual representation of how an observation $\psi(x_c)$ is interpreted as an assignment $\xi(\text{x'}_c)\$$ during scheduling.

We see that $\beta \subset \alpha$. Thus, if we receive an observation $\psi(x_c)$ earlier than $l^+(x_c)$, we may safely buffer by applying the execution strategy from an assignment of $\psi(x_c) = \xi(x_c') = l^+(x_c)$. □

Next, we address the case where $\psi(x_c) > u^-(x_c)$.

Lemma 12. *If a contingent event, $x_c \in X_c$, will be observed after the bounds of x_c', $\psi(x_c) > u^-(x_c)$, we imagine we have received it by assigning $\xi(x_c') = u^-(x_c)$ in S'.*

Proof. We apply the same argument to *imagining* late events. We now consider an observation at the upper bounds of x_c', $\psi(x_c) = \xi(x_c') = u^-(x_c)$. We then have a new α representing the range of the earliest and latest assignments to $\xi(x_c)$,

$$\alpha = u^-(x_c) - g(x_c)$$

$$= [u^-(x_c), u^-(x_c)] - [\bar{\gamma}^-(x_c), \bar{\gamma}^+(x_c)]$$

$$\alpha = [u - (\bar{\gamma}^+(x_c) - \bar{\gamma}^-(x_c)), u]$$

Once again, if S' is fixed-delay controllable, there must exist an execution strategy for $\xi(x_c') = u^-(x_c)$. It follows that we can apply this execution strategy when $\xi(x_c) \in \alpha$.

If we receive a late observation, $\psi(x_c) = u^-(x_c) + \epsilon$, we find that $\xi(x_c)$ must fall in the range of a new β, where

$$\beta = \left[(u^-(x_c) + \epsilon) - g(x_c)\right] \cap [l, u]$$

$$= \left[[u^-(x_c) + \epsilon, u^-(x_c) + \epsilon] - [\bar{\gamma}^-(x_c), \bar{\gamma}^+(x_c)]\right] \cap [l, u]$$

$$= [u - (\bar{\gamma}^+(x_c) - \bar{\gamma}^-(x_c)) + \epsilon, u + \epsilon] \cap [l, u]$$

$$\beta = [u - (\bar{\gamma}^+(x_c) - \bar{\gamma}^-(x_c)) + \epsilon, u]$$

We find that $\beta \subset \alpha$ again and can safely imagine that we received $\psi(x_c) = u^-(x_c)$. Of course, we need not wait to receive a late observation of x_c only to assign it to a time in the past. During execution, if we have not received $\psi(x_c)$ by $u^-(x_c)$, we imagine an observation arrived at $\psi(x_c) = u^-(x_c)$ and thus assign $\xi(x'_c) = u^-(x_c)$. We then ignore the real observation of x_c that we receive later. $\square$

We have addressed the key issue of reconciling observations from S with the dispatchable form from S'. We now present a dispatcher and wrapper algorithms on top of FAST-EX that combine to add robustness for variable observation delay.

5.3 Modified FAST-EX for Variable Observation Delay

We present an overview of the scheduling and dispatching algorithms below with explanations following.

While we made a careful distinction between x_c and x'_c in our discussion of scheduling, in our implementation it was important to be able to easily replace one with another when looking up values in hash-tables and lists. For instance, to implement Equation 5.1, we receive x_c but key the fixed-delay function on x'_c. Rather than adding an additional translation layer, we give each temporal event in S a unique name, all of which get copied to their equivalent events in S'. Hash-tables are keyed on event names, vastly simplifying lookups in the AllMax graph, delay function, and elsewhere.

Let x be a temporal event, $x \forall x \in X_c \cup X_e$.

Input: AllMax Graph G; fixed-delay function $\gamma(x'_c)$; Observation $\psi(x_c)$

Output: Updated AllMax Graph G

Initialization:

$\xi(\text{x'}_c) \leftarrow \psi(x_c) - \gamma(x'_c)$;

VDC-FAST-EX-Update:

for $l \in S'.contingentLinks()$ **do**

$\quad x_c \leftarrow l.endpoint()$;

$\quad a, b \leftarrow l.bounds()$;

$\quad$ **if** $\bar{\gamma}^+(x_c) == \infty \ or \ \bar{\gamma}^+(x_c) == \bar{\gamma}^-(x_c)$ **then**

$\quad\quad \gamma'(x_c) \leftarrow \bar{\gamma}^+(x_c)$;

$\quad$ **endif**

$\quad$ **else if** $b - a < \bar{\gamma}^+(x_c) - \bar{\gamma}^-(x_c)$ **then**

$\quad\quad \gamma'(x_c) \leftarrow \infty$;

$\quad$ **endif**

$\quad$ **else**

$\quad\quad l.setBounds(a + \bar{\gamma}^+(x_c), b + \bar{\gamma}^-(x_c))$;

$\quad\quad \gamma'(x_c) \leftarrow 0$;

$\quad\quad$ **for** $l' \in x_c.outgoingReqLinks()$ **do**

$\quad\quad\quad u, v \leftarrow l'.bounds()$;

$\quad\quad\quad l'.setBounds(u - \bar{\gamma}^-(x_c), v - \bar{\gamma}^+(x_c))$;

$\quad\quad$ **end**

$\quad\quad$ **for** $l' \in x_c.incomingReqLinks()$ **do**

$\quad\quad\quad u, v \leftarrow l'.bounds()$;

$\quad\quad\quad l'.setBounds(u + \bar{\gamma}^+(x_c), v + \bar{\gamma}^-(x_c))$;

$\quad\quad$ **end**

$\quad$ **endif**

end

`return` S', γ'

Algorithm 2: Algorithm for updating the AllMax graph when an observation arrives

5.3.1   Real vs No-op Events

The introduction of buffering and imagining events creates a new distinction between temporal events: there are events that need to be executed by the agent and there are those events that do not. We call these *real* and *no-op* ("no operation") events. Both contingent *and* requirement events may fall into either category. Below, we present our rationale for the distinction between real and no-op events, and how we modify real-time execution decisions accordingly.

To start, both buffered and imagined contingent events are no-ops. Both cases represent time-points that we use to update our dispatchable form to maintain consistency with S'.

Consider the process of normalization of an STNU [27]. While building the labeled distance graph during a dynamic controllabillity check, we rewrite contingent links such that their lower bounds are always 0. For instance, for a contingent event C and free event E, $C - E \in [l, u]$, during normalization we create a new requirement event, C', fixed at the lower bound of the contingent link, and then shift the bounds of the contingent link to start at 0 while maintaining the original range, $u - l$. This results in two constraints: $E - C' \in [l, l]$ and $C - C' \in [0, u - l]$ that still reflect the original contingent link's semantics.

To a scheduler, there is no distinction between the semantics of a real event, as modeled by a human planner writing an STNU for an agent to execute, and C', an artifact of checking controllability. Both are modeled in the AllMax distance graph forming the basis of RTED generation. However, an agent does not need to execute any task in the outside world to satisfy $E - C'$. We take a view that the only information our agent has about the timepoints it should execute comes from the input STNU. Thus, we need RTEDs to reflect the distinction between requirement events that are *real*, meaning the agent is responsible for taking some action to execute them, and those that are *no-ops*, or algorithmic by-products that require no operation. This distinction naturally leads to the following addendum to the definition of RTEDs.

Definition 27. Event-No-op Pair

An *Event-No-op Pair*, ϕ, is a two-tuple, $\langle X, \nu \rangle$, where:

- X is an event in $X_e \cup X_c$,

- ν is a boolean, where if True, the event is a no-op, else real

Definition 28. RTED with Operational Distinction

A *Real-Time Execution Decision with Operational Distinction* is a two-tuple $\langle t, \Phi \rangle$, where:

- t is a time with domain $\mathbb{R}$,

- Φ is a set of ϕ to be executed at time t,

For convenience and simplicity, and given the similarities between RTED and RTED with Operational Distinction, future references to RTEDs will always mean RTEDs with Operational Distinctions.

5.4 Dynamic Dispatching

A *dynamic dispatcher* (or just "dispatcher") is an interface layer with a two-fold responsibility: it triggers the execution of RTEDs in the outside world, and it relays observations from the outside

world about the execution of events to the scheduler.

5.5   Rescheduling

The goal of this method is to dispatch future events as soon as possible. Contingent events may arrive earlier than expected due to the information lost during the variable-delay to fixed-delay STNU transformation process. Without this method, we would always be forced to buffer observations to the expected start time from the fixed-delay STNU in order to guarantee controllability of the rest of the STNU. Here, we create a new variable-delay STNU reflecting the resolutions of uncertainty so far, namely converting the original variable-delay contingent link to a free event set to lower==upper bounds matching its actual execution time, then re-perform controllability checks. If controllable, we get a new schedule that removes the need to buffer this contingent event. If not controllable, we do nothing, buffer the ctg event as planned, and continue dispatching against the original schedule.

Assume we have enough time to perform rescheduling, eg. the margin between the resolution of the ctg event and the lower bound of when we were expecting it is greater than the time it takes to perform rescheduling

It's a generate-and-test approach

Chapter 6

Coordinating Multiple Agents in Extreme Environments

This is a chapter on distributed communication.

THIS PAGE INTENTIONALLY LEFT BLANK

Chapter 7

Coordinating Multiple Agents in Extreme Environments

This is a chapter on distributed communication.

THIS PAGE INTENTIONALLY LEFT BLANK

Chapter 8

Evaluation

This is a chapter on evaluating all this stuff.

THIS PAGE INTENTIONALLY LEFT BLANK

Chapter 9

Discussion and Future Work

This is a discussion of stuff to do later.

THIS PAGE INTENTIONALLY LEFT BLANK

Bibliography

[1] N. Bhargava, C. Muise, and B. C. Williams, Variable-delay controllability, in *IJCAI International Joint Conference on Artificial Intelligence*, 2018, vol. 2018-July, pp. 46604666. doi: 10.24963/ijcai.2018/648.

[2] M. J. Miller, Decision Support System Development For Human Extravehicular Activity, Georgia Institute of Technology, 2017.

[3] D. Wang and B. C. Williams, TBurton: A divide and conquer temporal planner, in *Proceedings of the National Conference on Artificial Intelligence*, 2015, vol. 5, pp. 34093417.

[4] J. W. McBarron, Past, present, and future: The U.S. EVA Program, *Acta astronautica*, vol. 32, no. 1, pp. 514, 1994, doi: 10.1016/0094-5765(94)90143-0.

[5] C. P. Sonnett, REPORT of the AD HOC WORKING GROUP ON APOLLO EXPERIMENTS AND TRAINING on the SCIENTIFIC ASPECTS OF THE APOLLO PROGRAM, NASA, 1963.

[6] J. M. Hurtado, K. Young, J. E. Bleacher, W. B. Garry, and J. W. Rice, Field geologic observation and sample collection strategies for planetary surface exploration: Insights from the 2010 Desert RATS geologist crewmembers, *Acta astronautica*, vol. 90, no. 2, pp. 344355, 2013, doi: 10.1016/j.actaastro.2011.10.015.

[7] K. Young and T. Graff, Planetary Science Context for EVA, in *NASA EVA Exploration Workshop*, 2020, p. 40.

[8] A. Kanelakos, Artemis EVA Flight Operations - Preparing for Lunar EVA Training & Execution, in *NASA EVA Exploration Workshop*, 2020, p. 47.

[9] C. Campbell, Advanced EMU Portable Life Support System (PLSS) and Shuttle/ISS EMU Schematics, a Comparison, *42nd international conference on environmental systems*, pp. 118, 2012, doi: 10.2514/6.2012-3411.

[10] E. S. Patterson and D. D. Woods, Shift changes, updates, and the on-call architecture in space shuttle mission control, *Computer supported cooperative work*, vol. 10, no. 3-4, p. 27, 2001, doi: 10.1023/A:1012705926828.

[11] S. J. Payler *et al.*, Developing Intra-EVA Science Support Team Practices for a Human Mission to Mars, *Astrobiology*, vol. 19, no. 3, pp. 387400, 2019, doi: 10.1089/ast.2018.1846.

[12] D. A. Coan, Exploration EVA System Concept of Operations Summary for Artemis Phase 1 Lunar Surface Mission, NASA, Houston, TX, 2020.

[13] M. A. Seibert, D. S. Lim, M. J. Miller, D. Santiago-Materese, and M. T. Downs, Developing Future Deep-Space Telecommunication Architectures: A Historical Look at the Benefits of

Analog Research on the Development of Solar System Internetworking for Future Human Spaceflight, *Astrobiology*, vol. 19, no. 3, pp. 462477, Mar. 2019, doi: 10.1089/ast.2018.1915.

[14] M. J. Miller and K. M. Feigh, *Addressing the envisioned world problem: A case study in human spaceflight operations*, vol. 5. Cambridge University Press, 2019. doi: 10.1017/dsj.2019.2.

[15] M. J. Miller, K. M. McGuire, and K. M. Feigh, Information flow model of human extravehicular activity operations, *Ieee aerospace conference proceedings*, vol. 2015-June, 2015, doi: 10.1109/AERO.2015.7118942.

[16] M. J. Miller, K. M. McGuire, and K. M. Feigh, Decision Support System Requirements Definition for Human Extravehicular Activity Based on Cognitive Work Analysis, *Journal of cognitive engineering and decision making*, vol. 11, no. 2, pp. 136165, 2017, doi: 10.1177/1555343416672112.

NO_ITEM_DATA:Sehlke2019

[18] B. C. Williams, M. D. Ingham, S. H. Chung, and P. H. Elliott, Model-based programming of intelligent embedded systems and robotic space explorers, *Proceedings of the ieee*, vol. 91, no. 1, pp. 212236, 2003, doi: 10.1109/JPROC.2002.805828.

[19] B. C. Williams and R. J. Ragno, Conflict-directed A* and its role in model-based embedded systems, *Discrete applied mathematics*, vol. 155, no. 12, pp. 15621595, 2007, doi: 10.1016/j.dam.2005.10.022.

[20] R. E. Fikes and N. J. Nilsson, Strips: A new approach to the application of theorem proving to problem solving, *Artificial intelligence*, vol. 2, no. 3-4, pp. 189208, 1971, doi: 10.1016/0004-3702(71)90010-5.

[21] R. Dechter, I. Meiri, and J. Pearl, Temporal constraint networks, *Artificial intelligence*, vol. 49, no. 1-3, pp. 6195, 1991, doi: 10.1016/0004-3702(91)90006-6.

[22] E. Fernández González, Generative Multi-Robot Task and Motion Planning Over Long Horizons, Massachusetts Institute of Technology, 2018.

[23] J. Chen, B. C. Williams, and C. Fan, Optimal mixed discrete-continuous planning for linear hybrid systems, in *HSCC 2021 - Proceedings of the 24th International Conference on Hybrid Systems: Computation and Control (part of CPS-IoT Week)*, 2021, vol. 1. doi: 10.1145/3447928.3456654.

[24] T. Vidal and H. Fargier, Handling Contingency in Temporal Constraint Networks: From Consistency to Controllabilities, *Journal of experimental and theoretical artificial intelligence*, vol. 11, no. 1, pp. 2345, 1999, doi: 10.1080/095281399146607.

[25] P. H. Morris, N. Muscettola, and T. Vidal, Dynamic Control Of Plans With Temporal Uncertainty, 2001.

[26] P. Morris and N. Muscettola, Temporal dynamic controllability revisited, *Proceedings of the national conference on artificial intelligence*, vol. 3, pp. 11931198, 2005.

[27] P. Morris, A structural characterization of temporal dynamic controllability, *Lecture notes in computer science (including subseries lecture notes in artificial intelligence and lecture notes in bioinformatics)*, vol. 4204 LNCS, pp. 375389, 2006, doi: 10.1007/11889205_28.

[28] P. Morris, Dynamic controllability and dispatchability relationships, *Lecture notes in computer science (including subseries lecture notes in artificial intelligence and lecture notes in bioinformatics)*, vol. 8451 LNCS, no. Dc, pp. 464479, 2014, doi: 10.1007/978-3-319-07046-9_33.

[29] A. Cimatti, A. Micheli, and M. Roveri, Solving temporal problems with uncertainty using SMT: Strong Controllability, *Constraints*, vol. 20, no. 1, pp. 129, 2012, doi: 10.1007/s10601-014-9167-5.

[30] M. Ingham, R. Ragno, A. Wehowsky, and B. Williams, The Reactive Model-based Programming Language, MIT Space Systems and Artificial Intelligence Laboratories, 2002.

[31] G. Kiczales, J. Des Rivières, and D. G. Bobrow, *The art of the metaobject protocol*. Cambridge, Mass: MIT Press, 1991.

[32] M. Fox and D. Long, PDDL2.1: An extension to PDDL for expressing temporal planning domains, *Journal of artificial intelligence research*, vol. 20, pp. 61124, 2003, doi: 10.1613/jair.1129.

[33] N. Bhargava, C. Muise, T. Vaquero, and B. Williams, Delay Controllability : Multi-Agent Coordination under Communication Delay, Massachusetts Institute of Technology, Cambridge, MA, 2018.

[34] L. Hunsberger, Efficient execution of dynamically controllable simple temporal networks with uncertainty, *Acta informatica*, vol. 53, no. 2, pp. 89147, 2016, doi: 10.1007/s00236-015-0227-0.

[35] N. Bhargava, Multi-Agent Coordination under Limited Communication, Massachusetts Institute of Technology, 2020.

[36] N. Bhargava, C. Muise, T. Vaquero, and B. Williams, Managing communication costs under temporal uncertainty, in *IJCAI International Joint Conference on Artificial Intelligence*, 2018, vol. 2018-July, pp. 8490. doi: 10.24963/ijcai.2018/12.

[37] L. Hunsberger, A faster execution algorithm for dynamically controllable STNUs, *Proceedings of the 20th international symposium on temporal representation and reasoning*, pp. 2633, 2013, doi: 10.1109/TIME.2013.13.

[38] L. Hunsberger, Fixing the semantics for dynamic controllability and providing a more practical characterization of dynamic execution strategies, *Time 2009 - 16th international symposium on temporal representation and reasoning*, pp. 155162, 2009, doi: 10.1109/TIME.2009.25.